

Diwali Sales Dataset

Dhanachote_W

2023-11-06

Introduction

Hi, my name is Dhanachote. I'm working on a mini-project to analyze and practice the R language, build my data visualization skills, and create a report using Rmarkdown. For this project, I'm using the 'Diwali Sales' dataset sourced from Kaggle, which you can find at the link below. Therefore, I cleaned the dataset in Excel before importing it into R programming.

Reference



Diwali Sales Dataset Source from Kaggle

Table

- Part I: Questions and Data Visualization
 - Question 1: What is the most popular product category during Diwali sales?
 - Question 2: Which customer group has the highest total purchase amount?
 - Question 3: Which state has the highest total customer purchase amount?
 - Question 4 : State / Zone

- * Question 4.1: Count the number of orders for each zone and calculate the average order size for each zone.
- * Question 4.2: Calculate the total number of orders across all zones.
- * Question 4.3: For each zone, identify the product category with the highest number of orders.
- * Question 4.4: Identify the product category with the highest total sales amount in the each zone.
- Question 5: Generation
 - * Question 5.1: Which customer has spent the most money? To which generation do they belong?
 - * Question 5.2: Calculate the average order amount for Generation Z customers.
 - * Question 5.3: Determine the number of male and female customers
 - * Question 5.4: Identify the number of customers belonging to each generation.
- Question 6: Identify the percentage amount for each occupation.
- Question 7: What is the total amount for each age group?
- Part II: Machine Learning

Part I Questions and Data Visualization

In this part, I asking a serveral questions to understanding in the dataset. I use library such as `tidyverse`, `readr`, `dplyr`, `ggplot2`, and `caret` before query the dataset.

Install packages and download library

```
## Library
library(readr)
library(dplyr)
library(tidyverse)
library(ggplot2)
library(caret)
```

```
## import read.csv()

diwali_sales <- tibble(read.csv("Diwali_Sales_Data_main.csv"))
```

Import dataset from .csv

```
## check NA

diwali_sales %>%
  complete.cases()%>%
  mean()
```

Check NA

```
## [1] 0
```

I use `complete.cases()` to check for missing values and then use `mean()` to calculate the logical mean (TRUE = 1, FALSE = 0). If the mean value is 0, it means that the dataset is complete and has no missing values.

Tip: You can check logical values by using the `as.logical()` function. For example, `as.logical(0)` evaluates to FALSE, while 1 evaluates to TRUE.

```
## select column

ds_cl <- diwali_sales %>%
  select(User_ID,
         Cust_name,
         Product_ID,
         Product_Category,
         Gender,
         Age,
         Marital_Status,
         State,
         Zone,
         Occupation,
         Orders,
         Amount) %>%
  drop_na()
```

Prepare the dataset

Question 1: What is the most popular product category during Diwali sales?

```
q1 <- diwali_sales %>%
  group_by(Product_Category) %>%
  summarise(Total_amount = sum(Orders * Amount)) %>%
  arrange(desc(Total_amount)) %>%
  head(5)

print(q1)
```

```
## # A tibble: 5 x 2
##   Product_Category      Total_amount
##   <chr>                <dbl>
## 1 Food                 83591272.
## 2 Clothing & Apparel   41164094
## 3 Electronics & Gadgets 39315276
## 4 Footwear & Shoes     38731504.
## 5 Furniture           13660130.
```

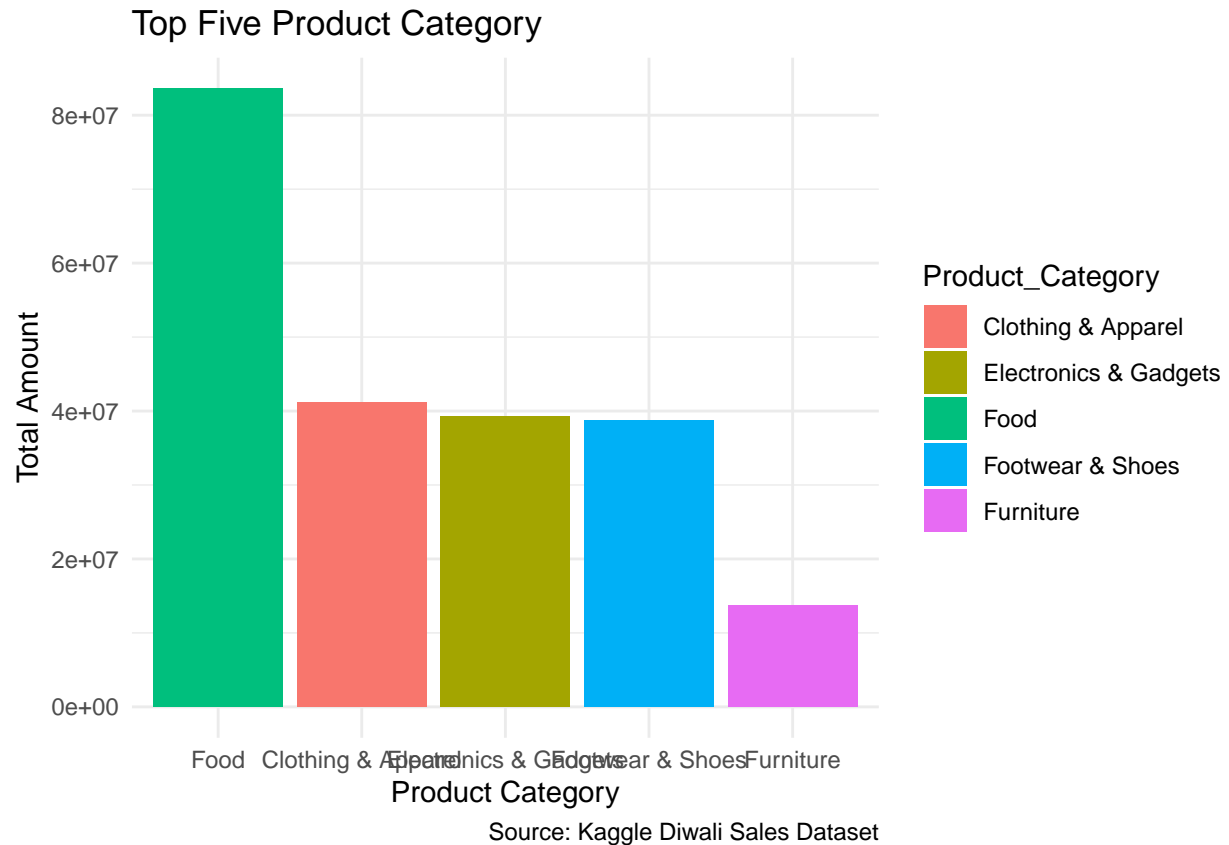
```
message("The most product category total amount is: ",
        head(q1$Product_Category,1),
        " ",
        "and",
        " ",
        "The total amount is: $ ",
        head(q1$Total_amount,1))
```

```
## The most product category total amount is: Food and The total amount is: $ 83591271.5
```

```
ds_cl %>%
  group_by(Product_Category) %>%
  summarise(
    total_amount = sum(Amount * Orders)
  ) %>%
  top_n(5) %>%
  arrange(desc(total_amount)) %>%
  ggplot(aes(reorder(Product_Category, -total_amount), total_amount, fill = Product_Category)) +
  geom_col() +
  theme_minimal() +
  labs(
    title = "Top Five Product Category",
    x = "Product Category",
    y = "Total Amount",
    caption = "Source: Kaggle Diwali Sales Dataset"
  )
```

Plot 1:

```
## Selecting by total_amount
```



Question 2: Which customer has the highest total purchase amount?

```
## Question2 : Who is the most spending amount
```

```
q2 <- diwali_sales %>%
  group_by(Cust_name) %>%
  summarise(
    Total_amount = sum(Orders * Amount)
  ) %>%
  arrange(desc(Total_amount)) %>%
  head(5)

print(q2)
```

```
## # A tibble: 5 x 2
##   Cust_name Total_amount
##   <chr>         <dbl>
## 1 Vishakha     972730
## 2 Alejandro    718053
## 3 Vasudev      698923
## 4 Sudevi       686455
## 5 Lalita       683635
```

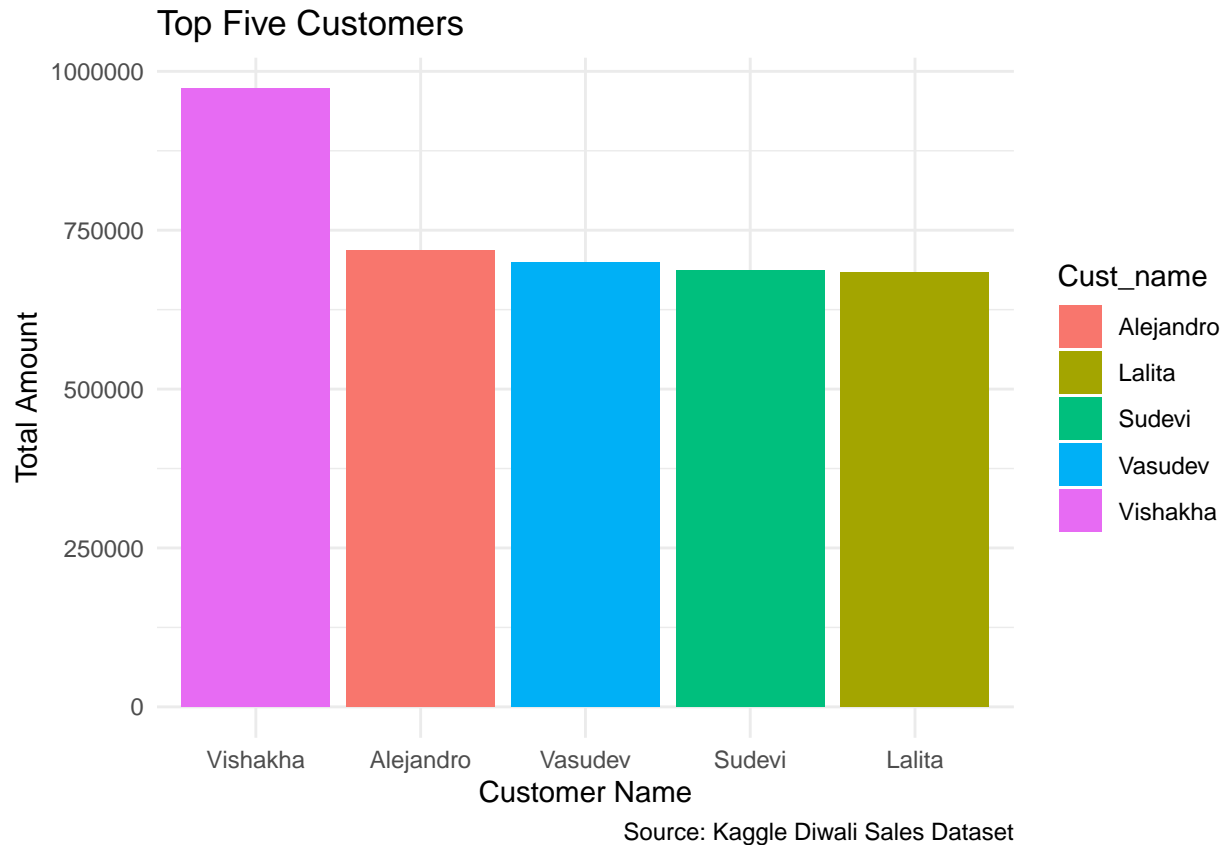
```
message(head(q2$Cust_name,1),
        " ",
        "is the most spending amount in $ ",
        head(q2$Total_amount,1),
        " product category")
```

```
## Vishakha is the most spending amount in $ 972730 product category
```

```
ds_cl %>%
  group_by(Cust_name) %>%
  summarise(
    total_amount = sum(Amount * Orders)
  ) %>%
  top_n(5) %>%
  ggplot(aes(reorder(Cust_name, -total_amount), total_amount, fill = Cust_name)) +
  geom_col() +
  theme_minimal() +
  labs(
    title = "Top Five Customers",
    x = "Customer Name",
    y = "Total Amount",
    caption = "Source: Kaggle Diwali Sales Dataset"
  )
```

Plot 2

```
## Selecting by total_amount
```



Question 3: Which state has the highest total customer purchase amount?

```
q3 <- diwali_sales%>%
  select(Cust_name,
         Gender,
         State,
         Zone,
         Orders,
         Amount) %>%
  mutate(Total_amount = Orders * Amount) %>%
  filter(Total_amount > 50000) %>%
  arrange(desc(Total_amount)) %>%
  head(10)

print(q3)
```

```
## # A tibble: 10 x 7
##   Cust_name Gender State      Zone    Orders Amount Total_amount
##   <chr>      <chr> <chr>    <chr>    <int>  <dbl>    <dbl>
## 1 Balk      F      Uttar Pradesh Central      4 23841     95364
## 2 Ginny     F      Andhra Pradesh Southern     4 23800     95200
```

##	3	Vasudev	M	Andhra Pradesh	Southern	4	23718	94872
##	4	Ellis	F	Andhra Pradesh	Southern	4	23546	94184
##	5	Mahima	F	Andhra Pradesh	Southern	4	23451	93804
##	6	Daniels	F	Andhra Pradesh	Southern	4	23302	93208
##	7	Mike	M	Himachal Pradesh	Northern	4	23267	93068
##	8	Dean	F	Andhra Pradesh	Southern	4	23252	93008
##	9	Zypern	M	Andhra Pradesh	Southern	4	23239	92956
##	10	Abhijit	F	Andhra Pradesh	Southern	4	23066	92264

```
message("Which state is the most user spending is: ",
        head(q3$State,1))
```


Which state is the most user spending is: Uttar Pradesh

Question 4 : State / Zone

In this question, I provided a chart that explains which state has the highest total amount before starting the questions.

Plot 4: Relationship with State and Amount

```
ds_cl %>%
  group_by(State) %>%
  ggplot(data = ds_cl,
        mapping = aes(x = Amount, y = State, fill = State)) +
  geom_col()
```

```
# count zone
diwali_sales %>%
  count(Zone)
```

Question 4.1: Count the number of orders for each zone and calculate the average order size for each zone.

```
## # A tibble: 5 x 2
##   Zone      n
##   <chr>  <int>
## 1 Central  4296
## 2 Eastern   814
## 3 Northern 1491
## 4 Southern 2695
## 5 Western  1955
```

```
avg_orders <- mean(diwali_sales$Orders)
```

```
## total and percentage avg orders
```

```
diwali_sales %>%
  select(Zone,
```

```

    Product_Category,
    Orders,
    Amount) %>%
filter(Orders >= avg_orders) %>%
group_by(Zone) %>%
summarise(total_avg_orders = n()) %>%
mutate(pct_avg_orders = c(total_avg_orders / sum(total_avg_orders)) * 100) %>%
arrange(desc(total_avg_orders))

```

```

## # A tibble: 5 x 3
##   Zone      total_avg_orders pct_avg_orders
##   <chr>          <int>          <dbl>
## 1 Central            2087            37.7
## 2 Southern           1347            24.3
## 3 Western            989            17.8
## 4 Northern            722            13.0
## 5 Eastern            398             7.18

```

```

# total orders
diwali_sales %>%
  select(Zone,
    Product_Category,
    Orders) %>%
  group_by(Zone) %>%
  summarise(total_orders = sum(Orders))

```

Question 4.2: Calculate the total number of orders across all zones.

```

## # A tibble: 5 x 2
##   Zone      total_orders
##   <chr>          <int>
## 1 Central            10640
## 2 Eastern             2015
## 3 Northern            3727
## 4 Southern            6744
## 5 Western             4881


```

Plot 4.2: Calculate the total number of orders across all zones.

```

ds_cl %>%
  group_by(Zone) %>%
  summarise(total_orders = sum(Orders)) %>%
  arrange(desc(total_orders)) %>%
  ggplot(mapping = aes(x = reorder(Zone, -total_orders), y = total_orders, fill = Zone)) +
  geom_col() +
  labs(
    title = "Total Orders by Zone",
    x = "Zone",
    y = "Total Orders",
    caption = "Source: Kaggle Diwali Sales Dataset"
  )

```

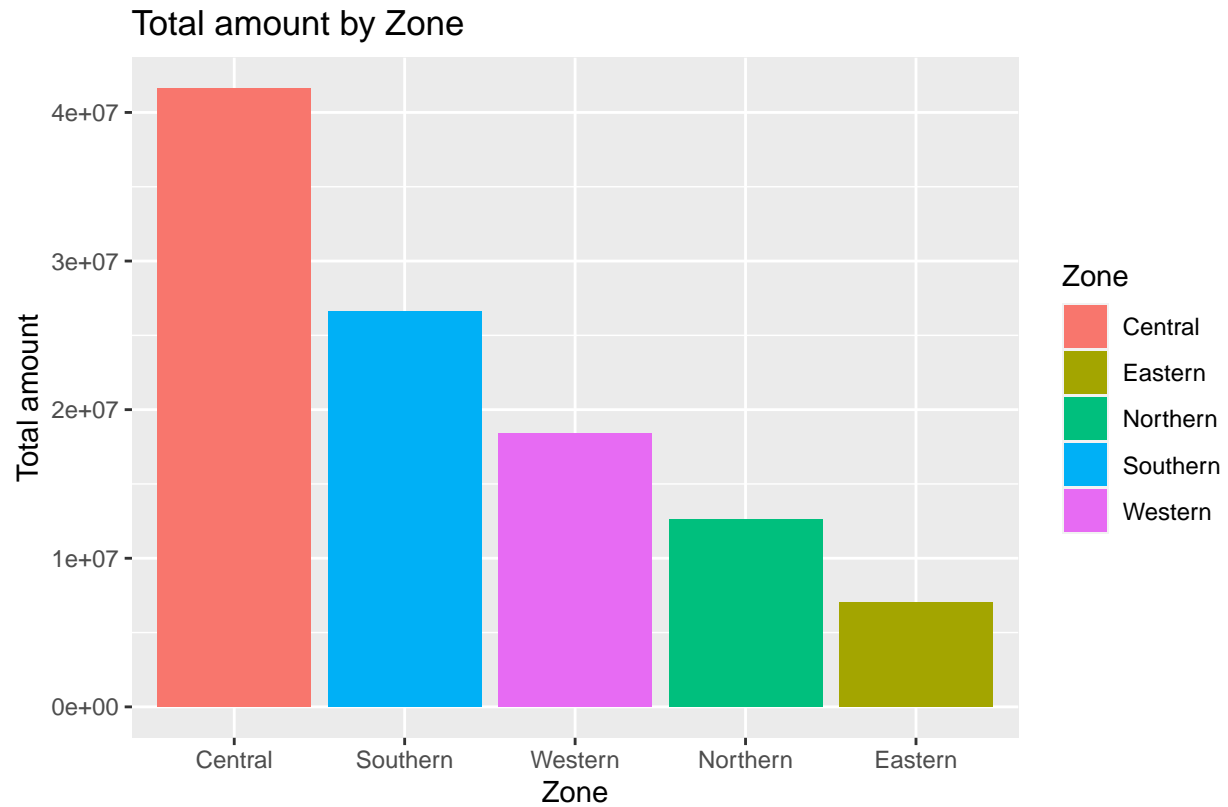


Question 4.3: For each zone, identify the product category with the highest number of orders.

The summary of this question involves visualizing the highest amount of orders by zone.

Plot 4.3: For each zone, identify the product category with the highest number of orders.

```
ds_cl %>%
  group_by(Zone) %>%
  summarise(total_amount = sum(Amount)) %>%
  arrange(desc(total_amount)) %>%
  ggplot(
    mapping = aes(x = reorder(Zone, -total_amount), y = total_amount, fill = Zone)) +
  geom_col() +
  labs(
    title = "Total amount by Zone",
    x = "Zone",
    y = "Total amount",
    caption = "Source: Kaggle Diwali Sales Dataset"
  )
```



Source: Kaggle Diwali Sales Dataset

Central

```
## Central

diwali_sales %>%
  select(Zone,
         Product_Category,
         Orders,
         Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Central") %>%
  summarize(Total_order = sum(Orders)) %>%
  arrange(desc(Total_order)) %>%
  head(2)
```

```
## # A tibble: 2 x 2
##   Product_Category Total_order
##   <chr>             <int>
## 1 Food              2452
## 2 Clothing & Apparel 2243
```

Western

```
## Western

diwali_sales %>%
```

```

select(Zone,
       Product_Category,
       Orders,
       Amount) %>%
group_by(Product_Category) %>%
filter(Zone == "Western") %>%
summarize(Total_order = sum(Orders)) %>%
arrange(desc(Total_order)) %>%
head(2)

```

```

## # A tibble: 2 x 2
##   Product_Category Total_order
##   <chr>             <int>
## 1 Food              1435
## 2 Clothing & Apparel 1118

```

Southern

```
## Southern
```

```

diwali_sales %>%
  select(Zone,
         Product_Category,
         Orders,
         Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Southern") %>%
  summarize(Total_order = sum(Orders)) %>%
  arrange(desc(Total_order)) %>%
  head(2)

```

```

## # A tibble: 2 x 2
##   Product_Category Total_order
##   <chr>             <int>
## 1 Clothing & Apparel 1449
## 2 Food              1313

```

Eastern

```
## Eastern
```

```

diwali_sales %>%
  select(Zone,
         Product_Category,
         Orders,
         Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Eastern") %>%
  summarize(Total_order = sum(Orders)) %>%
  arrange(desc(Total_order)) %>%
  head(2)

```

```
## # A tibble: 2 x 2
##   Product_Category      Total_order
##   <chr>                <int>
## 1 Clothing & Apparel      721
## 2 Electronics & Gadgets  552
```

Northern

```
## Northern
```

```
diwali_sales %>%
  select(Zone,
         Product_Category,
         Orders,
         Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Northern") %>%
  summarize(Total_order = sum(Orders)) %>%
  arrange(desc(Total_order)) %>%
  head(2)
```

```
## # A tibble: 2 x 2
##   Product_Category      Total_order
##   <chr>                <int>
## 1 Electronics & Gadgets  1316
## 2 Clothing & Apparel    1103
```

```
# which product category is the most spent in Central?
```

```
diwali_sales %>%
  select(Zone,
         Product_Category,
         Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Central") %>%
  summarize(Total_Amount = sum(Amount)) %>%
  arrange(desc(Total_Amount)) %>%
  head(10)
```

Question 4.4: Identify the product category with the highest total sales amount each zone.

```
## # A tibble: 10 x 2
##   Product_Category      Total_Amount
##   <chr>                <dbl>
## 1 Food                13685560.
## 2 Footwear & Shoes    8468991.
## 3 Clothing & Apparel  5744202
## 4 Electronics & Gadgets 3907342
## 5 Furniture          2067178
## 6 Beauty              1420386
```

```
## 7 Games & Toys          1382910
## 8 Stationery            1266360.
## 9 Sports Products      1023434
## 10 Decor                532002
```

which product category is the most spent in Western?

```
diwali_sales %>%
  select(Zone, Product_Category, Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Western") %>%
  summarize(Total_Amount = sum(Amount)) %>%
  arrange(desc(Total_Amount)) %>%
  head(10)
```

```
## # A tibble: 10 x 2
##   Product_Category Total_Amount
##   <chr>            <dbl>
## 1 Food             7764072
## 2 Clothing & Apparel 2698534
## 3 Electronics & Gadgets 2230270
## 4 Sports Products  1041627
## 5 Games & Toys      979280
## 6 Furniture        966517
## 7 Footwear & Shoes  929332
## 8 Household items   377449
## 9 Tupperware       362751
## 10 Stationery      246502
```

which product category is the most spent in Southern?

```
diwali_sales %>%
  select(Zone, Product_Category, Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Southern") %>%
  summarize(Total_Amount = sum(Amount)) %>%
  arrange(desc(Total_Amount)) %>%
  head(10)
```

```
## # A tibble: 10 x 2
##   Product_Category Total_Amount
##   <chr>            <dbl>
## 1 Food             7326461
## 2 Footwear & Shoes  5436504
## 3 Electronics & Gadgets 3882459
## 4 Clothing & Apparel 3649326
## 5 Furniture        1295949.
## 6 Auto             1237453.
## 7 Games & Toys      1001066
## 8 Sports Products   845462
## 9 Household items   526144
## 10 Books            478963
```

```
# which product category is the most spent in Eastern?

diwali_sales %>%
  select(Zone, Product_Category, Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Eastern") %>%
  summarize(Total_Amount = sum(Amount)) %>%
  arrange(desc(Total_Amount)) %>%
  head(10)
```

```
## # A tibble: 10 x 2
##   Product_Category      Total_Amount
##   <chr>                <dbl>
## 1 Food                 2209777
## 2 Clothing & Apparel   1772370
## 3 Electronics & Gadgets 1632873
## 4 Footwear & Shoes     315519
## 5 Games & Toys         303982
## 6 Furniture            302380
## 7 Sports Products      236266
## 8 Auto                 79632
## 9 Tupperware           54008
## 10 Household items     48448
```

```
# which product category is the most spent in Northern?

diwali_sales %>%
  select(Zone, Product_Category, Amount) %>%
  group_by(Product_Category) %>%
  filter(Zone == "Northern") %>%
  summarize(Total_Amount = sum(Amount)) %>%
  arrange(desc(Total_Amount)) %>%
  head(10)
```

```
## # A tibble: 10 x 2
##   Product_Category      Total_Amount
##   <chr>                <dbl>
## 1 Electronics & Gadgets 3990902
## 2 Food                 2948013
## 3 Clothing & Apparel   2630587
## 4 Furniture            808028
## 5 Games & Toys         664456
## 6 Sports Products      489144
## 7 Footwear & Shoes     424863
## 8 Books                140896
## 9 Tupperware           98820
## 10 Household items     94974
```

Question 5: Generation

Since the dataset does not have a Generation column, I need to create one.


```
## create generation table

diwali_sales_gen <- diwali_sales %>%
  select(User_ID,
         Cust_name,
         Product_ID,
         Gender,
         Age.Group,
         Age,
         Marital_Status,
         Marital_Status,
         State,
         Zone,
         Occupation,
         Product_Category,
         Orders,
         Amount) %>%
  mutate(Generation = ifelse(Age < 26, "Gen Z",
                             ifelse(Age < 42, "Millennials",
                                     ifelse(Age < 58, "Gen X",
                                             ifelse(Age < 68, "Baby Boomer", "Silent Generation")))))
```

Question 5.1: Which customer has spent the most money? To which generation do they belong?

```
diwali_sales_gen %>%
  select(User_ID,
         Cust_name,
         Gender,
         Age,
         Orders,
         Amount,
         Occupation,
         Generation) %>%
  mutate(total_amount = Orders * Amount) %>%
  arrange(desc(total_amount)) %>%
  head(10)
```

Question 5.1: Which customer has spent the most money? To which generation do they belong?

```
## # A tibble: 10 x 9
##   User_ID Cust_name Gender  Age Orders Amount Occupation Generation
##   <int> <chr>   <chr> <int> <int> <dbl> <chr>   <chr>
## 1 1001132 Balk     F      25     4 23841 Lawyer   Gen Z
## 2 1003650 Ginny    F      26     4 23800. Media   Millennials
## 3 1001680 Vasudev  M      26     4 23718 Automobile Millennials
## 4 1000113 Ellis    F      19     4 23546 Govt     Gen Z
## 5 1004736 Mahima   F      25     4 23451 Banking Gen Z
## 6 1004505 Daniels  F      55     4 23302 Healthcare Gen X
## 7 1002520 Mike     M      72     4 23267 Media    Silent Generati~
```

```
## 8 1003111 Dean      F      25      4 23252 Banking      Gen Z
## 9 1001182 Zypern    M      16      4 23239 Food Processing Gen Z
## 10 1001726 Abhijit  F      32      4 23066 Retail        Millennials
## # i 1 more variable: total_amount <dbl>
```

```
## Avg amount gen z

diwali_sales_gen %>%
  group_by(Generation) %>%
  filter(Generation == "Gen Z") %>%
  summarize(
    total_amount = sum(Amount),
    avg_amount = mean(Amount),
    percentage_amount = (sum(Amount) / sum(diwali_sales_gen$Amount)) * 100
  )
```

Question 5.2: Calculate the average order amount for Generation Z customers.

```
## # A tibble: 1 x 4
##   Generation total_amount avg_amount percentage_amount
##   <chr>         <dbl>      <dbl>          <dbl>
## 1 Gen Z          19940385      9168.           18.8
```

```
## How many Male and Female in Gender?

diwali_sales %>%
  count(Gender)
```

Question 5.3: Determine the number of male and female customers

```
## # A tibble: 2 x 2
##   Gender      n
##   <chr> <int>
## 1 F      7842
## 2 M      3409
```

```
## Plot 3: Which gender are the most come to shopping at store

ggplot(data = ds_cl,
  mapping = aes(x = Gender, fill = Gender)) +
  geom_bar() +
  theme_minimal() +
  labs(
    title = "The most gender shopping at store",
    x = "Gender",
    y = "Count",
    caption = "Source: Kaggle Diwali Sales Dataset"
  )
```



```
## How many Generation in Customer?

diwali_sales_gen %>%
  count(Generation)
```

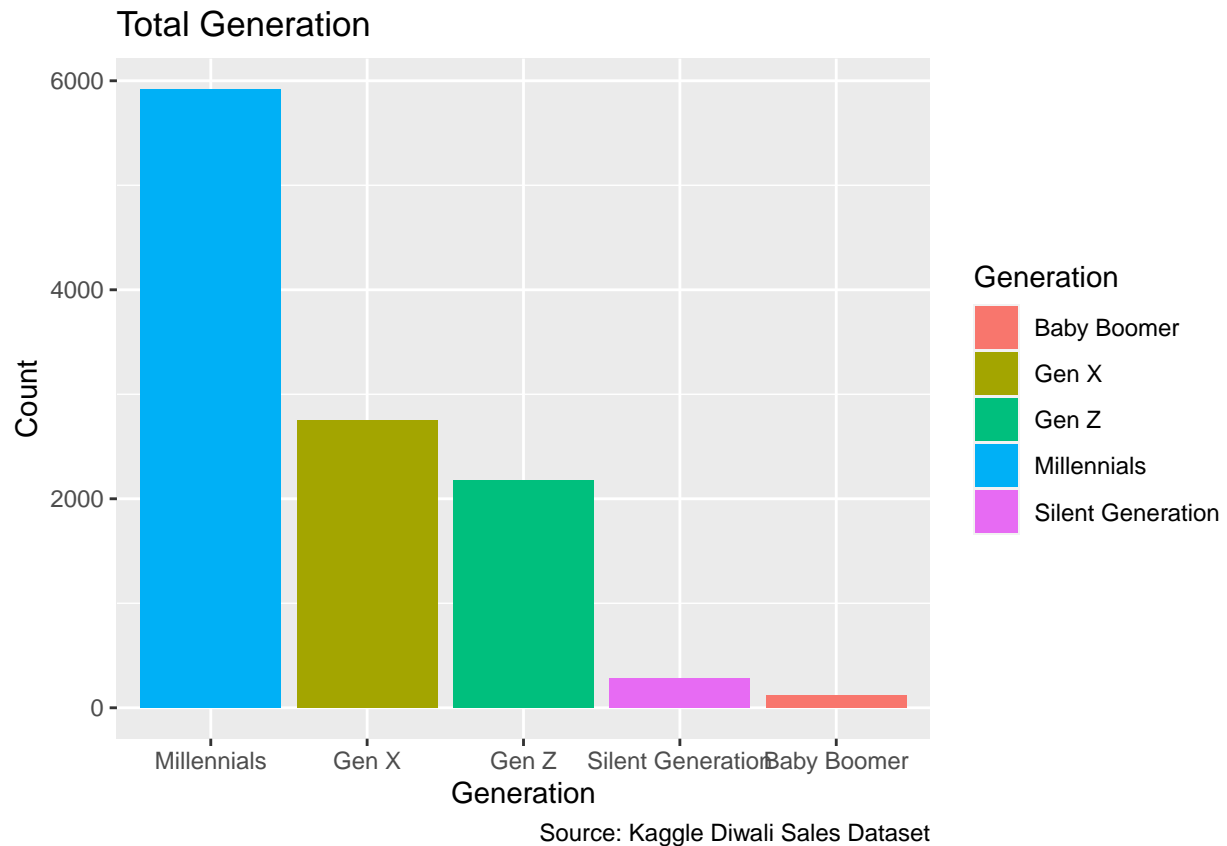
Question 5.4: Identify the number of customers belonging to each generation.

```
## # A tibble: 5 x 2
##   Generation      n
##   <chr>         <int>
## 1 Baby Boomer    123
## 2 Gen X         2754
## 3 Gen Z         2175
## 4 Millennials   5916
## 5 Silent Generation 283
```

```
## Plot 5: Total Generation

diwali_sales_gen %>%
  group_by(Generation) %>%
  count() %>%
  ggplot(mapping = aes(x = reorder(Generation, -n), y = n, fill = Generation)) +
  geom_col() +
```

```
labs(
  title = "Total Generation",
  x = "Generation",
  y = "Count",
  caption = "Source: Kaggle Diwali Sales Dataset"
)
```



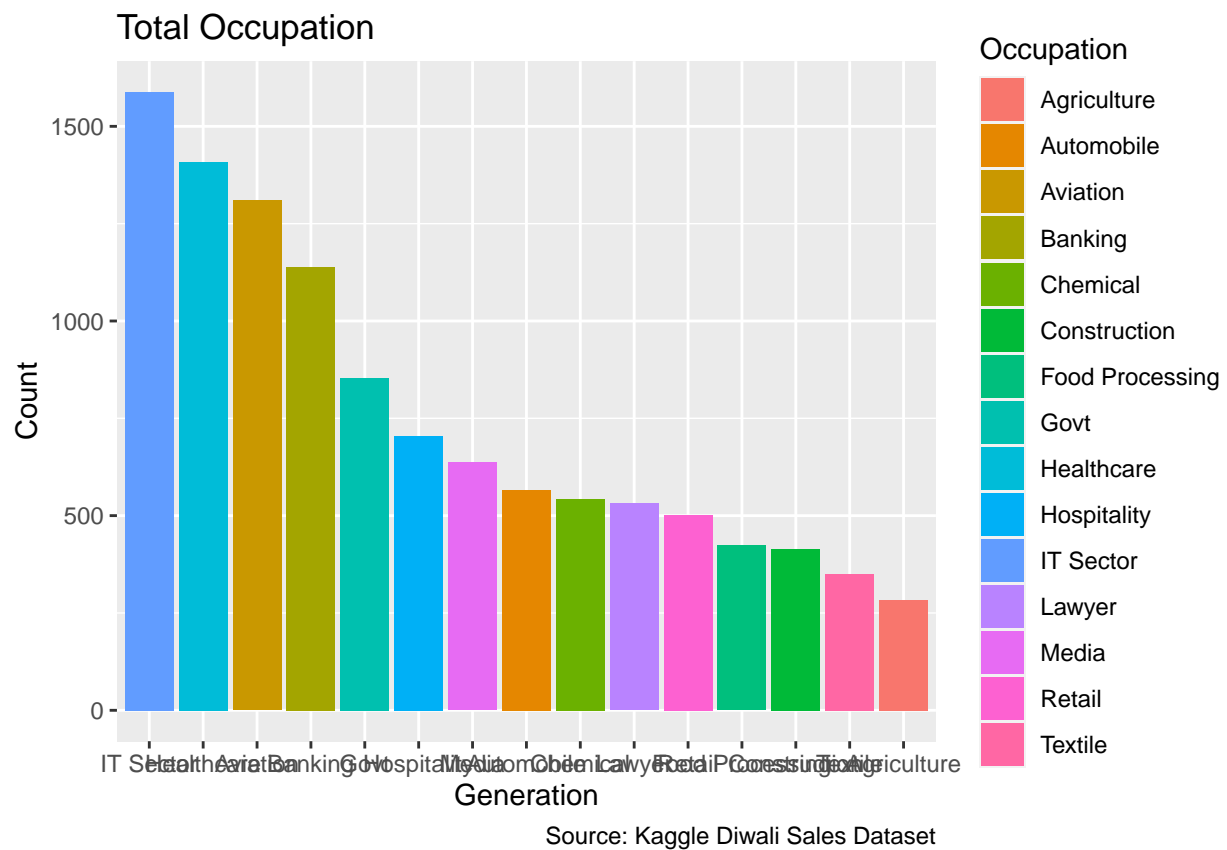
Occupationa

Before moving on to the next question, I provided the total occupation count from the dataset.

Plot 6: Total Occupation

```
diwali_sales_gen %>%
  group_by(Occupation) %>%
  count() %>%
  ggplot(mapping = aes(x = reorder(Occupation, -n), y = n, fill = Occupation))+
  geom_col() +
  labs(
    title = "Total Occupation",
    x = "Generation",
    y = "Count",
```

```
caption = "Source: Kaggle Diwali Sales Dataset"
)
```



Question 6: Identify the percentage amount for each occupation.

```
## Question 6: Identify the percentage amount for each occupation.
```

```
avg_amount <- mean(diwali_sales$Amount)
```

```
diwali_sales %>%
  select(Occupation,
         Orders,
         Amount) %>%
  group_by(Occupation) %>%
  summarise(total_avg_amount = n(),
            pct_amount = (sum(Amount) / sum(diwali_sales_gen$Amount)) * 100) %>%
  arrange(desc(total_avg_amount))
```

```
## # A tibble: 15 x 3
##   Occupation      total_avg_amount pct_amount
##   <chr>          <int>      <dbl>
## 1 IT Sector      1588      13.9
## 2 Healthcare     1408      12.3
```

## 3 Aviation	1310	11.9
## 4 Banking	1139	10.1
## 5 Govt	854	8.02
## 6 Hospitality	705	6.00
## 7 Media	637	5.93
## 8 Automobile	566	5.05
## 9 Chemical	542	4.99
## 10 Lawyer	531	4.69
## 11 Retail	501	4.50
## 12 Food Processing	423	3.83
## 13 Construction	414	3.39
## 14 Textile	350	3.02
## 15 Agriculture	283	2.44

Age Group

Question 7: What is the total amount for each age group?

```
## Question 7: What is the total amount for each age group?
```


```
diwali_sales %>%
  select(Age.Group,
         Product_Category,
         Amount,
         Orders) %>%
  group_by(Age.Group) %>%
  summarise(total_amount = sum(Amount),
            pct_amount = (sum(Amount) / sum(diwali_sales_gen$Amount)) * 100) %>%
  arrange(desc(total_amount))
```

```
## # A tibble: 7 x 3
##   Age.Group total_amount pct_amount
##   <chr>         <dbl>         <dbl>
## 1 26-35      42613444.         40.1
## 2 36-45      22144995.         20.8
## 3 18-25      17240732          16.2
## 4 46-50       9207844           8.67
## 5 51-55       8261477           7.78
## 6 55+        4080987           3.84
## 7 0-17        2699653           2.54
```

```
## Plot 7: What is the total amount for each age group?
```

```
diwali_sales %>%
  group_by(Age.Group) %>%
  summarise(total_amount = sum(Amount)) %>%
  arrange(desc(total_amount)) %>%
  ggplot(mapping = aes( x= reorder(Age.Group, -total_amount), y = total_amount, fill = Age.Group)) +
  geom_col() +
```

```
labs(
  title = "Total amount by Age group",
  x = "Age groups",
  y = "Total Amount",
  caption = "Source: Kaggle Diwali Sales Datasets"
)
```



Summary

In summary, the most prominent product category in the Diwali Sales dataset is food. The dataset shows that in the state of Uttar Pradesh, there are a lot of customers buying during Diwali. In terms of zones, the Central zone has the highest total amount. Lastly, the most common customer age range is 26-35, representing Generation Z. By the way, I've noticed that females spend more than males in this dataset.

Part II: Machine Learning

What is Machine Learning

Machine learning is a type of artificial intelligence (AI) focused on building computer systems that learn from data.

Before delving into machine learning. I am using machine learning to predict the amount of sales during Diwali. In this section, I have broken down the topic to make it understandable for everyone, demonstrating how I code by following these steps:

- Install package
 - Import dataset
 - Prepare data
 - Split data
 - Train & Test data
 - Scoring
 - Evaluate model
-

```
## install.packages
install.packages("readr")
install.packages("dplyr")
install.packages("tidyverse")
install.packages("ggplot2")

## Library
library(readr)
library(dplyr)
library(tidyverse)
library(ggplot2)
library(caret)

## import read.csv()

diwali_sales <- tibble(read.csv("Diwali_Sales_Data_main.csv"))

## select column

ds_cl <- diwali_sales %>%
  select(User_ID,
         Cust_name,
         Product_ID,
         Product_Category,
         Gender,
         Age,
         Marital_Status,
         State,
         Zone,
         Occupation,
         Orders,
         Amount) %>%
  drop_na()

## check NA
```




```
ds_cl %>%  
  complete.cases() %>%  
  mean()
```

Install package and Import dataset

```
## [1] 1
```

```
ggplot(ds_cl, aes(Amount)) +  
  geom_density() +  
  theme_minimal() +  
  labs(  
    title = "Quick Visualization of Amount",  
    x = "Amount",  
    caption = "Source: Kaggle Diwali Sales Datasets"  
  )
```



Quick Visualization of Amount

```
## split data

sp_data <- function(ds_cl, train_size = 0.7) {
  set.seed(42)
  n <- nrow(diwali_sales)
  id <- sample(n, size = n * train_size)
  tr_df <- ds_cl[id,]
  ts_df <- ds_cl[-id,]
  list(train = tr_df,
        test = ts_df)
}

## prep data

prep_data <- sp_data(ds_cl)
tr_df <- prep_data[[1]]
ts_df <- prep_data[[2]]
```

Split and Prep Data

```
## train model

set.seed(42)
lm_model <- train(Amount ~ Orders ,
                  data = tr_df,
                  method = "lm")

## view model

lm_model
```

Train data

```
## Linear Regression
##
## 7875 samples
##    1 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 7875, 7875, 7875, 7875, 7875, 7875, ...
## Resampling results:
##
##    RMSE      Rsquared    MAE
##    5225.28  0.0005284117  4312.628
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
# score <- prediction

p <- predict(lm_model, newdata = ts_df)
```

Scoring

```
## evaluate
# mae mse rmse
## create function

cal_mae <- function(actual, pred) {
  error <- actual - pred
  return(mean(abs(error)))
}

cal_mse <- function(actual, pred) {
  error <- actual - pred
  mean(error ** 2)
}

cal_rmse <- function(actual, pred) {
  error <- actual - pred
  sqrt(mean(error ** 2))
}

# check result

cal_mae(ts_df$Amount, p)
```

Evaluate

```
## [1] 4272.277
```

```
cal_mse(ts_df$Amount, p)
```

```
## [1] 26999848
```

```
cal_rmse(ts_df$Amount, p)
```

```
## [1] 5196.138
```

```
## Summary

lm_model$finalModel %>%
  summary()
```

Summary

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9597  -3992  -1337   3247  14498
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9676.97    144.54   66.949  <2e-16 ***
## Orders       -80.27     53.05   -1.513    0.13
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5243 on 7873 degrees of freedom
## Multiple R-squared:  0.0002907, Adjusted R-squared:  0.0001637
## F-statistic: 2.289 on 1 and 7873 DF,  p-value: 0.1303
```